

Homework 2

Problem 1

- a. This data was provided to build a predictive model for the rating variable. The first step was to clear the environment and load the MASS and leaps libraries. I downloaded and stored the dataset locally, and since it is a CSV file, I next used the read.csv command to load the file and examined it a bit.

```
> names(dats)
[1] "name" "mfr" "type" "calories" "protein" "fat" "sodium" "fiber"
[9] "carbo" "sugars" "potass" "vitamins" "shelf" "weight" "cups" "rating"
> head(dats)
      name mfr type calories protein fat sodium fiber carbo sugars potass
1 100% Bran N C 70 4 1 130 10.0 5.0 6 280
2 100% Natural Bran Q C 120 3 5 15 2.0 8.0 8 135
3 All-Bran K C 70 4 1 260 9.0 7.0 5 320
4 All-Bran with Extra Fiber K C 50 4 0 140 14.0 8.0 0 330
5 Almond Delight R C 110 2 2 200 1.0 14.0 8 -1
6 Apple Cinnamon Cheerios G C 110 2 2 180 1.5 10.5 10 70
      vitamins shelf weight cups rating
1 25 3 1 0.33 68.40297
2 0 3 1 1.00 33.98368
3 25 3 1 0.33 59.42551
4 25 3 1 0.50 93.70491
5 25 3 1 0.75 34.38484
6 25 1 1 0.75 29.50954
```

I divided the data next, with a 70% training and 30% testing ratio, setting a seed of 123 such that my results can be replicated. Since the mfr and type variables are categorical, and since there isn't much variance in type, I have removed the first the 3 columns. I next fit a simple linear model to rating, and then used the predict function. To calculate the mean square error, I took the mean of the square of the residuals component. My result is 6.6207e-14.

```
> print(MSE)
[1] 6.620702e-14
> |
```

- b. For part b, I performed forward subset selection. The default value for method in the regsubsets function is exhaustive, as mentioned in the lab, therefore I gave the value forward to the method parameter. A summary of the selection shows the best subset selections with different number of variables.

```
weight cups FALSE FALSE
1 subsets of each size up to 12
Selection Algorithm: forward
calories protein fat sodium fiber carbo sugars potass vitamins shelf weight cups
1 (1) 0 0 0 0 0 0 0 0 0 0 0 0
2 (1) 0 0 0 0 0 0 0 0 0 0 0 0
3 (1) 0 0 0 0 0 0 0 0 0 0 0 0
4 (1) 0 0 0 0 0 0 0 0 0 0 0 0
5 (1) 0 0 0 0 0 0 0 0 0 0 0 0
6 (1) 0 0 0 0 0 0 0 0 0 0 0 0
7 (1) 0 0 0 0 0 0 0 0 0 0 0 0
8 (1) 0 0 0 0 0 0 0 0 0 0 0 0
9 (1) 0 0 0 0 0 0 0 0 0 0 0 0
10 (1) 0 0 0 0 0 0 0 0 0 0 0 0
11 (1) 0 0 0 0 0 0 0 0 0 0 0 0
12 (1) 0 0 0 0 0 0 0 0 0 0 0 0
```

I experimented a bit to see if changing the number of best subsets for each size revealed anything unexpected, but it didn't, so I let the best single subset of each size stay.

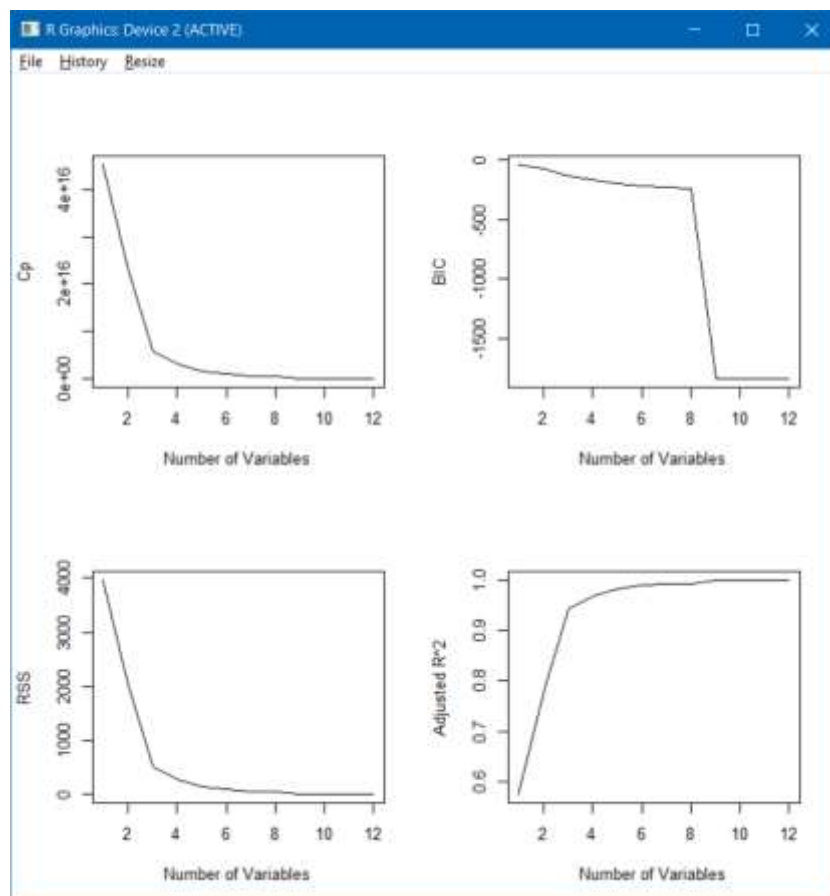
c. For part c as well I did the same as for part b.

```

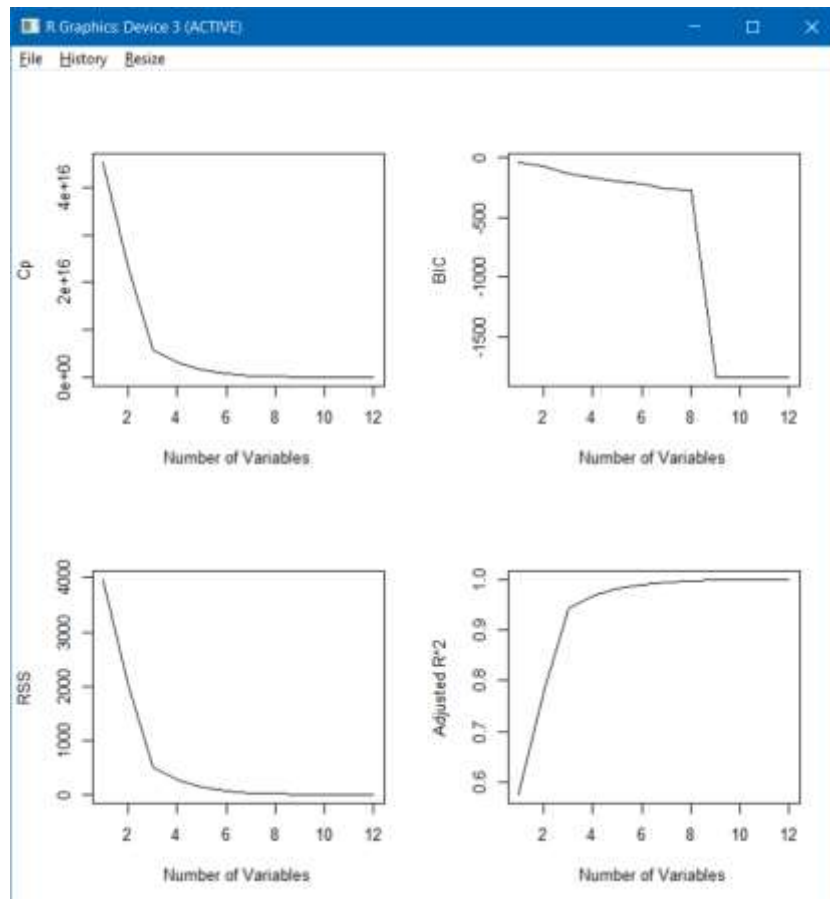
> stepAIC <- stepAIC(
+   weight ~ 1, FALSE, FALSE,
+   cups ~ 1, FALSE, FALSE,
+   1 subsets of each size up to 12,
+   Selection = exhaustive)
> stepAIC
calories protein fat sodium fiber carbo sugars potass vitamins shelf weight cups
1 (1) 11 11 11 11 11 11 11 11 11 11 11 11
2 (1) 11 11 11 11 11 11 11 11 11 11 11 11
3 (1) 11 11 11 11 11 11 11 11 11 11 11 11
4 (1) 11 11 11 11 11 11 11 11 11 11 11 11
5 (1) 11 11 11 11 11 11 11 11 11 11 11 11
6 (1) 11 11 11 11 11 11 11 11 11 11 11 11
7 (1) 11 11 11 11 11 11 11 11 11 11 11 11
8 (1) 11 11 11 11 11 11 11 11 11 11 11 11
9 (1) 11 11 11 11 11 11 11 11 11 11 11 11
10 (1) 11 11 11 11 11 11 11 11 11 11 11 11
11 (1) 11 11 11 11 11 11 11 11 11 11 11 11
12 (1) 11 11 11 11 11 11 11 11 11 11 11 11
> I

```

In addition to the subset selections, I also displayed the Cp, BIC, RSS and adjusted R squared for both forward and exhaustive.



Forward



Exhaustive

As is clear from the graphs, the forward and exhaustive selections are quite similar. There is a sharp decline in BIC for both forward and exhaustive subset selections around the 9 variable mark. Based on my results, I would say that the exhaustive subset selection is the best model.

Problem 2

- The first step was to load the train and test data. Since only 2s and 3s are to be focused on for this problem, I filtered the data to only include them. I also added all the k values in a list.
- For linear regression, I used the V1 variable as the focus for fitting a line to the data. To check its prediction accuracy, I ran a loop through the values of the predictions and essentially rounded off the values such that the prediction is binary, either 2 or 3. Then, I simply mathematically calculated the error and stored it in a variable. The error was roughly 4.12%

```
> error  
[1] 0.04120879  
> |
```

- c. For knn classification, as suggested in the pseudo code, I ran a loop through the list of k values, using the knn algorithm function each time and mathematically calculating the error each time and appending it to the store_error list. The testing errors corresponding to ascending order of k values (1, 3, 5, 7, 9, 11, 13, 15):

```
> store_error
[1] 0.02472527 0.03021978 0.03021978 0.03021978 0.03571429 0.03571429
[7] 0.03296703 0.03846154
```

k-value	Error
1	2.47%
3	3.02%
5	3.02%
7	3.02%
9	3.57%
11	3.57%
13	3.29%
15	3.85%

It appears that knn classification is better for all cases than linear regression. The error percentage seems to increase with increase in k-value except for a decline in error between 11 and 13. Beyond 15 as well, the error follows an increasing trend.

```
> store_error2
[1] 0.000000000 0.004319654 0.005759539 0.005759539 0.007919366 0.007919366 0.007919366 0.009359251
```

The training error for each of the k values is as above